

Reviewer 1

Summary

This paper presents Question Maker, an instructor-facing web application for AI-assisted assessment authoring. The system combines RAG with structured workflows. The paper positions itself against general-purpose LLM tools by emphasizing workflow-level features such as enforced instructor review before deployment, explicit variant tracking, and course-scoped generation. An author-led validation using a three-question midterm in a computer organization course is reported.

Strengths

The problem is well-motivated. Assessment authoring is genuinely time-intensive, and the gap between standalone LLM generation and deployable, workflow-integrated assessment tooling is real. The paper clearly articulates why general-purpose LLM tools are insufficient and maps its design objectives to specific system features.

The design decision to enforce review gating at the persistence layer rather than only at export is a useful insight. The paper honestly reports what did not work, including topic misclassification during variant generation and overly restrictive filtering with sparse question banks. This kind of candid reflection is valuable for an experience report.

Weaknesses

The most significant weakness is that the paper reads more like a system design document than an experience report. The reader needs to understand concretely what it looks and feels like to use this tool, and the paper does not provide that properly. There is insufficient information about the actual interactions, example questions or variants shown, screenshots of what the instructor sees during generation or review, and concrete illustration of what a RAG-grounded prompt or response looks like in practice. Figure 2 shows the assessment assembly interface but is not walked through in the text. A reader cannot form a mental model of using this system from the paper as written.

The "Initial Validation" (Section 6) is extremely thin. The entire evaluation is author-led validation on a single three-question assessment, conducted prior to any deployment. The paper explicitly states it does not represent observed instructor adoption or classroom use. For a long-form experience report, this is quite enough. It should provide a rich reflection on what did or didn't work during actual use. There is no formative data, no instructor feedback, and no indication of whether the system

was ever used by anyone other than the authors, so there is insufficient information about the experience using the tool.

In terms of writing, the paper's description of the system is heavily abstract and repetitive. Many sections restate the same points about review gating, persistent variants, and topic constraints using nearly identical phrasing. For instance, the concept of "review gating" is described in the abstract, several times in substantially similar terms each time. The writing could be more concise and better organized.

The font in Figure 1 is too small to be informative. Overall, there should be more figures shown in this paper. For a tool paper, the visual presentation of the tool's interface is critical. More and larger annotated screenshots showing the key workflow stages would substantially improve the paper.

Score

2 - Borderline

In your opinion, does this submission make a contribution to the WCCCE community?

No

Would you advocate for acceptance of this submission?

No

Reviewer 2

The paper introduces QuestionMaker, a generative AI tool for generating and organizing variant or new questions, and organizing them into an assignment.

Strengths:

- Organizing and tagging question variants with metadata is useful
- Requiring instructor review before questions are added to the pool is sensible
- Condition question generation on Bloom's taxonomy is an interesting and potentially useful idea

Weaknesses:

- Organization of paper is sometimes confusing; E.g., reasoning level is mentioned in passing on page 3, and only explained to be the level of Bloom's hierarchy on page 4. Similarly "topics" are explained in section 4.4, but mentioned earlier without explanation
- The initial validation seems weak; it is purely qualitative, and the variant tested seems extremely easy (narrowing the scope of an existing problem), creating "a multiple-choice variant that focused on a subset of the original question's outcome while using the same underlying code and execution context."; it would have been nice to see the resulting question. Even if self-assessment is the only way forwards, I would have liked to know what proportion of the questions were usable, or how much time was required in refining the genAI's output (even a small sample of 10 examples)

would be better than a single example with a takeaway of "variant generation worked"). Similarly, I would like to see more details in 6.5; the paper states that "Some drafts required substantial editing" -- how much time was required? What was the comparison with starting from scratch? How many times was the system tested? Of "correct ambiguity, adjust difficulty, or better align with intended learning objectives" how frequent were they? How diverse were the outputs? Did the system respect the requested level of Bloom's hierarchy? Without a better sense of where the system is or isn't working, I am very hesitant to accept this paper. It is good to know that the filtering works, and that requiring instructor review and approval for exportation works, but to me the real unknown here is whether LLMs are sufficiently reliable to actually be useful instructional partners.

Concluding: I think the design is interesting, specialized for teaching, and potentially quite useful. However, I would really have liked to see more detail in the evaluation. Even if this is an "early evaluation" at least reporting the number of attempts, the rate at which the reported limitations occurred, and the average amount of human time in correcting the outputs (vs. writing from scratch) would go a very long way to convincing me that this tool is genuinely promising (or at the very least, highlighting where there is still room for future improvement and development). I strongly encourage the authors to continue this interesting avenue of work, and expand the evaluation in future.

Score

3 - Meaningful

In your opinion, does this submission make a contribution to the WCCCE community?

Yes

Would you advocate for acceptance of this submission?

No

Reviewer 3

The paper introduces Question Maker which is a new tool to help instructors use AI to generate and manage assessment questions. The paper explains the advantages of using such a system over a plain LLM's chat interface to be: providing course context through RAG using course content, tracking how AI generated question variants were derived from base questions, and tracking/managing what questions have been reviewed for a human.

The strength of the paper is in how the tool, if proven through wider tests, could help instructors manage assessment creation in an integrated workflow. The tool would provide help creating new questions based on course content, new question variants based on existing questions, and on managing how questions are assembled into complete assessments.

There are a number of areas where the paper could use improvement. First, the paper's review of the system's utility seemed to be confined to a single three question assessment for a single course. Furthermore, the assessment in question seems not to have been used in an actual offering of the course. Therefore, the general usability of the tool has not yet been established in spite of the tool's design ideas being strong.

The paper would be strengthened with the inclusion of practical examples which showed what the tool could do given some course content and sample questions. It would help to see what novel questions, and what variants of existing questions, the system could create. This discussion could ground some of the discussion in section 6 which was somewhat abstract. Discussion of how the system manages question difficulty could likewise be added to this.

Minor points:

- I found the paper repeated the same ideas numerous times. Examples include AI generated content only being stored after human review, question variants are linked to base questions, and nothing can be exported unless reviewed by a human. I felt that the paper is too long for its novelty, impact, and depth of evaluation.

- Numerous phrases seemed to be more complex than needed. For example I would encourage simplifying or explaining phrases such as:

"rather than as ephemeral generation outputs."

"design emphasizes explicit workflow stages such as review gating and controlled persistence to ensure that instructor intervention is not optional or implicit but structurally enforced." - What is "controlled persistence"? "review gating" likewise seems a complicated way of saying it.

"These systems highlight the importance of structured representations of questions and metadata in deployed instructional contexts." - what are instructional contexts?

"foregrounds pre-deployment assessment workflows" - what does it mean to "foreground"?

- Referring to human-in-the-loop as "theoretical grounding" does not seem like the best choice of wording unless the paper wants to lean into the theoretical side of these frameworks.

- Section 2.4 mentions "Many prototypes acknowledge the need for human oversight..."; a citation here could help.

- Section 3.1: unclear what "enforced context isolation" is in this context.

- Figure 1's font is quite small compared to article font size (it is small at 300% viewing). The figure in general did not seem to strengthen my understanding of the system.

- Figure 2 is quite small with font viewable at 300%.

Score

2 - Borderline

In your opinion, does this submission make a contribution to the WCCCE community?

Yes

Would you advocate for acceptance of this submission?

No